



TMDB Movie Data Analysis and Building a Movie Recommendation System

- # Part 1: Data Cleaning/Transformation #### In this project, we will analyze movie data from the TMDB database, and try to extract some meaningful insight from it with Exploratory Data Analysis, visualization, and DataFrame manipulations. #### There are four parts to this end-to-end project: #### • Data download, data cleaning, feature engineering. #### • Exploratory Data Analysis. #### • Creating the final dataset and creating the recommendation algorithm. #### • Creating a Streamlit app to be deployed on Heroku. #### The first part of this project will focus on cleaning and manipulating the DataFrame, extracting any valuable information and removing any data that is not needed. #### The end result will be a movie recommendation system based on movie vector similarity.

- **id:** The ID of the movie (clear/unique identifier).
- **title:** The Official Title of the movie.
- **tagline:** The tagline of the movie.
- **release_date:** Theatrical Release Date of the movie.
- **genres:** Genres associated with the movie.
- **belongs_to_collection:** Gives information on the movie series/franchise the particular film belongs to.
- **original_language:** The language in which the movie was originally shot in.
- **budget_musd:** The budget of the movie in million dollars.
- **revenue_musd:** The total revenue of the movie in million dollars.
- **production_companies:** Production companies involved with the making of the movie.
- **production_countries:** Countries where the movie was shot/produced in.
- **vote_count:** The number of votes by users, as counted by TMDB.
- **vote_average:** The average rating of the movie.
- **popularity:** The Popularity Score assigned by TMDB.
- **runtime:** The runtime of the movie in minutes.
- **overview:** A brief blurb of the movie.
- **spoken_languages:** Spoken languages in the film.
- **poster_path:** The URL of the poster image.
- **cast:** (Main) Actors appearing in the movie.
- **cast_size:** number of Actors appearing in the movie.
- **director:** Director of the movie.
- **crew_size:** Size of the film crew (incl. director, excl. actors).

Loading the main libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from ast import literal_eval
import urllib
import requests
import json
import gzip
import time
from tqdm import tqdm

pd.options.display.max_columns = 20
```

Getting all valid movie IDs from TMDB's Daily File Exports

```
In [3]: def get_movie_ids():
        json_gz_path = 'http://files.tmbd.org/p/exports/movie_ids_10_28_2021.json.'
```

```

urllib.request.urlretrieve(json_gz_path, 'movie_ids_10_28_2021.json.gz') #
movie_ids = []
with gzip.open('movie_ids_10_28_2021.json.gz', 'r') as f:
    for line in f:
        movie_ids.append(json.loads(line)['id'])
return movie_ids

```

```
In [4]: movie_ids = get_movie_ids()
```

```
In [5]: print(f"First 10 movie ids: {movie_ids[:10]}")
print(f"# of movie ids: {format(len(movie_ids), ',d')}")
```

```

First 10 movie ids: [3924, 6124, 8773, 25449, 31975, 2, 3, 5, 6, 8]
# of movie ids: 650,449

```

Downloading the movie dataset using TMDB's API

```
In [18]: session = requests.Session()
session.headers = {
    "User-Agent": "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.3",
    "Accept-Encoding": "*",
    "Connection": "keep-alive"
}
```

```
In [22]: def retry(func, retries=10):
def retry_wrapper(*args, **kwargs):
    attempts = 0
    while attempts < retries:
        try:
            return func(*args, **kwargs)
        except requests.exceptions.RequestException as e:
            print(e)
            time.sleep(5)
            attempts += 1

    return retry_wrapper
```

```
In [30]: from api_key import api # Personal API key used to download the data.

failed_ids = []

@retry
def download_movie_data(movie_id, session=session):
    link = f'https://api.themoviedb.org/3/movie/{movie_id}?api_key={api}&language=en-US'
    response = session.get(link)
    if response.status_code == 200:
        return response.json()
    else:
        failed_ids.append(movie_id)
        return np.nan
```

```
In [ ]: movies = []
pbar = tqdm(total=len(movie_ids))
for movie_id in movie_ids:
    movies.append(download_movie_data(movie_id))
    pbar.update(1)
pbar.close()
```

```
In [117... len(movies)
```

```
Out[117]: 650643
```

```
In [118... len(failed_ids)
```

```
Out[118]: 194
```

Retrying failed movie requests

```
In [107... failed_movies = []  
for failed_movie_id in failed_ids:  
    failed_movies.append(download_movie_data(movie_id))
```

```
In [110... movies.extend(failed_movies)
```

```
In [111... len(movies)
```

```
Out[111]: 650643
```

Saving the resultant movies dataset

```
In [112... movies_series = pd.Series(movies).dropna()
```

```
In [113... df = pd.DataFrame.from_records(movies_series)
```

```
In [119... df.head()
```



```
In [132... credits_series = pd.Series(credits).dropna()
```

```
In [133... credits_df = pd.DataFrame(credits_series)
credits_df.to_csv('credits.csv', index=False)
del credits_df
```

Out[133]:

0

```
0 {'id': 3924, 'cast': [{'adult': False, 'gender': ...
1 {'id': 6124, 'cast': [{'adult': False, 'gender': ...
2 {'id': 8773, 'cast': [{'adult': False, 'gender': ...
3 {'id': 25449, 'cast': [{'adult': False, 'gende...
4 {'id': 31975, 'cast': [], 'crew': []}
5 {'id': 2, 'cast': [{'adult': False, 'gender': ...
6 {'id': 3, 'cast': [{'adult': False, 'gender': ...
7 {'id': 5, 'cast': [{'adult': False, 'gender': ...
8 {'id': 6, 'cast': [{'adult': False, 'gender': ...
9 {'id': 8, 'cast': [], 'crew': [{'adult': False...
10 {'id': 9, 'cast': [{'adult': False, 'gender': ...
11 {'id': 11, 'cast': [{'adult': False, 'gender':...
12 {'id': 12, 'cast': [{'adult': False, 'gender':...
13 {'id': 13, 'cast': [{'adult': False, 'gender':...
14 {'id': 14, 'cast': [{'adult': False, 'gender':...
15 {'id': 15, 'cast': [{'adult': False, 'gender':...
16 {'id': 16, 'cast': [{'adult': False, 'gender':...
17 {'id': 17, 'cast': [{'adult': False, 'gender':...
18 {'id': 18, 'cast': [{'adult': False, 'gender':...
19 {'id': 19, 'cast': [{'adult': False, 'gender':...
20 {'id': 20, 'cast': [{'adult': False, 'gender':...
21 {'id': 21, 'cast': [{'adult': False, 'gender':...
22 {'id': 22, 'cast': [{'adult': False, 'gender':...
23 {'id': 24, 'cast': [{'adult': False, 'gender':...
24 {'id': 25, 'cast': [{'adult': False, 'gender':...
25 {'id': 26, 'cast': [{'adult': False, 'gender':...
26 {'id': 27, 'cast': [{'adult': False, 'gender':...
27 {'id': 28, 'cast': [{'adult': False, 'gender':...
28 {'id': 33, 'cast': [{'adult': False, 'gender':...
29 {'id': 35, 'cast': [{'adult': False, 'gender':...
30 {'id': 38, 'cast': [{'adult': False, 'gender':...
31 {'id': 55, 'cast': [{'adult': False, 'gender':...
32 {'id': 58, 'cast': [{'adult': False, 'gender':...
33 {'id': 59, 'cast': [{'adult': False, 'gender':...
34 {'id': 62, 'cast': [{'adult': False, 'gender':...
35 {'id': 63, 'cast': [{'adult': False, 'gender':...
```

0

```
36  {'id': 64, 'cast': [{'adult': False, 'gender':...
37  {'id': 65, 'cast': [{'adult': False, 'gender':...
38  {'id': 66, 'cast': [{'adult': False, 'gender':...
39  {'id': 67, 'cast': [{'adult': False, 'gender':...
40  {'id': 68, 'cast': [{'adult': False, 'gender':...
41  {'id': 69, 'cast': [{'adult': False, 'gender':...
42  {'id': 70, 'cast': [{'adult': False, 'gender':...
43  {'id': 71, 'cast': [{'adult': False, 'gender':...
44  {'id': 73, 'cast': [{'adult': False, 'gender':...
45  {'id': 74, 'cast': [{'adult': False, 'gender':...
46  {'id': 75, 'cast': [{'adult': False, 'gender':...
47  {'id': 76, 'cast': [{'adult': False, 'gender':...
48  {'id': 77, 'cast': [{'adult': False, 'gender':...
49  {'id': 78, 'cast': [{'adult': False, 'gender':...
```

- All of our data are downloaded and ready to view.

Data Cleaning


Some of the data attributes are in json format, so we will have to extract the columns/variables from json dictionary strings.

```
In [2]: df = pd.read_csv('movies_metadata.csv', low_memory=False)
df
```


Out[2]:

	adult	backdrop_path	belongs_to_collection	budget	genres
0	False	/dvQj1GBZAZirz1skEEZyWH2ZqQP.jpg	{'id': 177062, 'name': 'Blondie Collection', '...	0	[{'id': 35, 'name': 'Comedy'}]
1	False	NaN	NaN	0	[{'id': 12, 'name': 'Adventure'}]
2	False	/uJlc4aNPF3Y8yAqahJTKBwgwPVW.jpg	NaN	0	[{'id': 18, 'name': 'Drama'}, {'id': 10749, 'n...
3	False	NaN	{'id': 425256, 'name': 'New World Disorder', '...	0	[]
4	False	NaN	NaN	0	[{'id': 10751, 'name': 'Family'}]
...
650444	False	NaN	NaN	0	[{'id': 99, 'name': 'Documentary'}]
650445	False	NaN	NaN	0	[{'id': 99, 'name': 'Documentary'}]
650446	False	NaN	NaN	0	[{'id': 99, 'name': 'Documentary'}]
650447	False	NaN	NaN	0	[{'id': 99, 'name': 'Documentary'}]
650448	False	NaN	NaN	0	[{'id': 99, 'name': 'Documentary'}]

650449 rows × 25 columns



In [3]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 650449 entries, 0 to 650448
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   adult                                650449 non-null  bool
1   backdrop_path                       166759 non-null  object
2   belongs_to_collection               16176 non-null   object
3   budget                             650449 non-null  int64
4   genres                             650449 non-null  object
5   homepage                           75382 non-null   object
6   id                                  650449 non-null  int64
7   imdb_id                            432724 non-null  object
8   original_language                  650449 non-null  object
9   original_title                     650446 non-null  object
10  overview                           538518 non-null  object
11  popularity                         650449 non-null  float64
12  poster_path                       438992 non-null  object
13  production_companies               650449 non-null  object
14  production_countries               650449 non-null  object
15  release_date                      590521 non-null  object
16  revenue                           650449 non-null  int64
17  runtime                           597471 non-null  float64
18  spoken_languages                   650449 non-null  object
19  status                             650449 non-null  object
20  tagline                           91811 non-null   object
21  title                             650446 non-null  object
22  video                             650449 non-null  bool
23  vote_average                      650449 non-null  float64
24  vote_count                        650449 non-null  int64
dtypes: bool(2), float64(3), int64(4), object(16)
memory usage: 115.4+ MB

```

Dropping irrelevant columns that will not be needed.

```

In [4]: df.drop(columns=['adult', 'imdb_id', 'original_title', 'video', 'homepage', 'b
df.head()

```

Out[4]:

	belongs_to_collection	budget	genres	id	original_language	overview	popularity
0	{'id': 177062, 'name': 'Blondie Collection', '...	0	[{'id': 35, 'name': 'Comedy'}]	3924	en	Blondie and Dagwood are about to celebrate the...	2.445 /o6
1	NaN	0	[{'id': 12, 'name': 'Adventure'}]	6124	de	Der Mann ohne Namen is a German adventure movi...	0.600
2	NaN	0	[{'id': 18, 'name': 'Drama'}, {'id': 10749, 'n...	8773	fr	Love at Twenty unites five directors from five...	4.985 /ε
3	{'id': 425256, 'name': 'New World Disorder', '...	0	[]	25449	en	Gee Atherton ripping the Worlds course the day...	1.337 /okQ\
4	NaN	0	[{'id': 10751, 'name': 'Family'}]	31975	en	Elmo is making a very, very super special surp...	0.600 /qK\

How to Handle Stringified JSON Columns

Columns that contain dictionaries or lists are treated as strings in pandas.

In order to access their entries, they will be converted to their appropriate form using the **ast** module.

```
In [5]: import json
import ast
```

```
In [6]: json_col = ['belongs_to_collection', 'genres', 'production_countries', 'produc
```

```
In [7]: df[json_col].head()
```

Out[7]:

	belongs_to_collection	genres	production_countries	production_companies	spoken_languages
0	{'id': 177062, 'name': 'Blondie Collection', '...	[{'id': 35, 'name': 'Comedy'}]	[{'iso_3166_1': 'US', 'name': 'United States of America'}]	[{'id': 5, 'logo_path': '/71BqEFAF4V3qjjMPCpLu...'}]	[{'english_name': 'English', 'iso_639_1': 'en'}]
1	NaN	[{'id': 12, 'name': 'Adventure'}]	[{'iso_3166_1': 'DE', 'name': 'Germany'}]	[]	[]
2	NaN	[{'id': 18, 'name': 'Drama'}, {'id': 10749, 'name': 'Family'}]	[{'iso_3166_1': 'DE', 'name': 'Germany'}, {'iso_3166_1': 'US', 'name': 'United States of America'}]	[{'id': 38936, 'logo_path': '/ypvTqUeQOxORhFEF...'}]	[{'english_name': 'German', 'iso_639_1': 'de'}]
3	{'id': 425256, 'name': 'New World Disorder', '...	[]	[]	[]	[{'english_name': 'English', 'iso_639_1': 'en'}]
4	NaN	[{'id': 10751, 'name': 'Family'}]	[]	[]	[]

```
In [2]: def lit_eval(column):
        return column.apply(lambda x: literal_eval(x) if isinstance(x, str) else None)
```

```
In [9]: for column in json_col:
        print(f"\n{'='*40}\nColumn Name: {column}\nOriginal dtype: {type(df[column])}\nNew dtype: {type(df[column].apply(lit_eval))}")
```

```
=====
Column Name: belongs_to_collection
Original dtype: <class 'str'>
New dtype: <class 'list'>
```

```
=====
Column Name: genres
Original dtype: <class 'str'>
New dtype: <class 'list'>
```

```
=====
Column Name: production_countries
Original dtype: <class 'str'>
New dtype: <class 'list'>
```

```
=====
Column Name: production_companies
Original dtype: <class 'str'>
New dtype: <class 'list'>
```

```
=====
Column Name: spoken_languages
Original dtype: <class 'str'>
New dtype: <class 'list'>
```

The end result is that all of the JSON columns that were not accessible due to being treated as strings are now in the form of

either a list or a dictionary, and are ready to be accessed/have their values extracted.

In [10]: `df.head()`

Out[10]:

	belongs_to_collection	budget	genres	id	original_language	overview	popularity	
0	{'id': 177062, 'name': 'Blondie Collection', '...	0	[{'id': 35, 'name': 'Comedy'}]	3924	en	Blondie and Dagwood are about to celebrate the...	2.445	/o6
1	NaN	0	[{'id': 12, 'name': 'Adventure'}]	6124	de	Der Mann ohne Namen is a German adventure movi...	0.600	
2	NaN	0	[{'id': 18, 'name': 'Drama'}, {'id': 10749, 'n...	8773	fr	Love at Twenty unites five directors from five...	4.985	/k
3	{'id': 425256, 'name': 'New World Disorder', '...	0	[]	25449	en	Gee Atherton ripping the Worlds course the day...	1.337	/okQ\
4	NaN	0	[{'id': 10751, 'name': 'Family'}]	31975	en	Elmo is making a very, very super special surp...	0.600	/qK\

Extracting the useful information from each JSON column.

Each one of the JSON columns has a "name" variable that is an attribute of each column. (ex. "Romance" under "genres", "en" under "spoken_languages", etc.)

This for loop checks the data type of each column. If it's in dictionary format, there is going to be only one "name" key for that row/column.

If it's in a dictionary format, then there might be multiple "name" keys for each dictionary in that list, which will be combined using a pipe ("|") symbol where applicable.

```
In [11]: for column in json_col:
          if type(df[column][0]) == list:
              df[column] = df[column].apply(lambda x: '|'.join(i['name'] for i in x))
          else:
              df[column] = df[column].apply(lambda x: x['name'] if isinstance(x, dict) else x)
```

Here is the end result for a specific row and each column that was converted. (Compare to same code above)

```
In [12]: for column in json_col:
          print(f'Column "{column}": {df[column][2]}\n{"="*100}')
```

Column "belongs_to_collection": nan
=====

Column "genres": Drama|Romance
=====

Column "production_countries": Germany|France|Italy|Japan|Poland
=====

Column "production_companies": Ulysse Productions|Unitec Films|Cinesecolo|Toho Co.|Beta Film|Film Polski|Zespół Filmowy "Kamera"
=====

Column "spoken_languages": Deutsch|Français|Italiano|日本語|Polski
=====

Replacing blank entries with NaN

Some of the rows have empty strings for each column attribute.

```
In [13]: for column in json_col:
          print(f"*** {column} ***")
          print(f"{df[column].value_counts(dropna=False).head()}\n{'='*50}")
```

```

*** belongs_to_collection ***
NaN                                634273
Our Gang: The Roach/MGM talkies      80
Hopalong Cassidy Collection         66
Our Gang: The Roach/Pathé silents   66
The Durango Kid Collection          65
Name: belongs_to_collection, dtype: int64
=====
*** genres ***
215618
Documentary      79318
Drama            67732
Comedy           41816
Music            16955
Name: genres, dtype: int64
=====
*** production_countries ***
266260
United States of America  95773
Germany                  23767
United Kingdom           21261
France                   20550
Name: production_countries, dtype: int64
=====
*** production_companies ***
371837
Metro-Goldwyn-Mayer      1970
Columbia Pictures        1758
Warner Bros. Pictures     1670
BBC                      1539
Name: production_companies, dtype: int64
=====
*** spoken_languages ***
272572
English      129138
Español       23604
Français    21461
日本語       20586
Name: spoken_languages, dtype: int64
=====

```

Using a simple for loop, each empty string will be replaced with NaN. The final result is below.

```

In [16]: for column in json_col:
          df[column].replace('', np.nan, inplace=True)

```

```

In [17]: for column in json_col:
          print(f"*** {column} ***")
          print(f"{df[column].value_counts(dropna=False).head()}\n{' '*50}")

```

```

*** belongs_to_collection ***
NaN                634273
Our Gang: The Roach/MGM talkies    80
Hopalong Cassidy Collection      66
Our Gang: The Roach/Pathé silents  66
The Durango Kid Collection        65
Name: belongs_to_collection, dtype: int64
=====
*** genres ***
NaN                215618
Documentary        79318
Drama              67732
Comedy             41816
Music              16955
Name: genres, dtype: int64
=====
*** production_countries ***
NaN                266260
United States of America    95773
Germany                   23767
United Kingdom            21261
France                    20550
Name: production_countries, dtype: int64
=====
*** production_companies ***
NaN                371837
Metro-Goldwyn-Mayer        1970
Columbia Pictures          1758
Warner Bros. Pictures       1670
BBC                        1539
Name: production_companies, dtype: int64
=====
*** spoken_languages ***
NaN                272572
English                 129138
Español                 23604
Français              21461
日本語                  20586
Name: spoken_languages, dtype: int64
=====

```

Replacing all zero values for "budget" and "revenue" with NaN

Changing the scale of the data to millions USD

```

In [18]: for column in ['budget', 'revenue']:
          df[column] = df[column].replace(0, np.nan)
          df[column] = df[column].div(1000000)
          df.rename(columns={column: column + '_musd'}, inplace=True)

```

Replacing all zero values for runtime with NaN


```
In [19]: print(f"Original:\n{df.runtime.value_counts(dropna=False).head()}\n{' '*50}")
df.runtime.replace(0, np.nan, inplace=True)
print(f"New:\n{df.runtime.value_counts(dropna=False).head()}\n{' '*50}")
```

Original:

```
0.0      111613
NaN       52978
90.0      21644
10.0       9944
7.0        9143
```

Name: runtime, dtype: int64

=====

New:

```
NaN      164591
90.0      21644
10.0       9944
7.0        9143
6.0        8845
```

Name: runtime, dtype: int64

=====

Replacing entries where vote_average equals zero with NaN

```
In [20]: df.loc[df.vote_count == 0, 'vote_average'] = np.nan
```

Converting "release_date" column to datetime format

```
In [21]: df['release_date'] = pd.to_datetime(df['release_date'])
df['release_date'].dtype
```

```
Out[21]: dtype('<M8[ns]')
```

Creating a "year" column containing the release year

```
In [22]: df['year'] = df['release_date'].apply(lambda x: str(x).split('-')[0] if len(st
```

```
In [23]: df.year.value_counts()
```

```
Out[23]: NaT      59928
2019      28403
2020      27592
2018      25765
2017      24644
...
1879         1
2030         1
1882         1
2050         1
2032         1
Name: year, Length: 156, dtype: int64
```

Replacing entries in "overview" column where the entry explicates missing overview

```
In [24]: df.overview.value_counts(dropna=False).head(30)
```

```
Out[24]: NaN
111931
No overview found.
1063
Mexican feature film
952
No overview found
309
No synopsis
303
No Overview
287
Plot unknown.
210
Testimonies of seven prominent photographers active in the juncture of the stu
dent movement of 1968
195
Know what this is about?
171
Documentary film.
169
No Synopsis
143
What the movie has in store for you, wait and watch this space for more update
s.
139
Coming Soon
131
Overview Coming Soon...
124
Feature film.
97
No description
93
Short film.
89
Mexican movie
87
Lift your butt, sculpt your core, and get an incredible total body transformat
ion with 80 unique "real-time" workouts and a Timed-Nutrition Eating Plan.
82

77
Transform :20 is a high-intensity, six-week workout and nutrition program that
will help transform your body and mind in just 20 minutes a day.
67
No overview.
62
»
59
Add the plot.
55
coming soon..
54
Japanese comedy film.
53
An Eternalism film.
52
shaw production
49
```

A short animation by Shintaro Kago.

49

Combat Bulletin was a new reel type series keeping folks up-to-date with the combat of World War II.

49

Name: overview, dtype: int64

```
In [25]: recurring_text = ['No overview found.', 'No overview found', 'No synopsis', 'N  
                        'Overview Coming Soon...', 'Feature film.', 'No description'  
                        'No description.', 'not available']  
  
df.replace(dict.fromkeys(recurring_text, np.nan), inplace=True)
```

Replacing all blank values in the dataset with NaN

```
In [26]: df.replace(dict.fromkeys([' ', ' ', ' ', ' '], np.nan), inplace=True)
```

Removing duplicate entries from the dataset

```
In [27]: print(f"Duplicate Rows:\n{df.duplicated(keep=False).value_counts()}")
```

```
Duplicate Rows:  
False      650254  
True         195  
dtype: int64
```

```
In [28]: df.drop_duplicates(inplace=True)
```

```
In [29]: print(f"Duplicate Rows:\n{df.duplicated(keep=False).value_counts()}")
```

```
Duplicate Rows:  
False      650255  
dtype: int64
```

Handling Missing Values & Removing Observations

```
In [30]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 650255 entries, 0 to 650254
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   belongs_to_collection                 16176 non-null  object
1   budget_musd                          31064 non-null  float64
2   genres                               434637 non-null object
3   id                                    650255 non-null int64
4   original_language                    650255 non-null object
5   overview                             535077 non-null object
6   popularity                           650255 non-null float64
7   poster_path                          438798 non-null object
8   production_companies                 278418 non-null object
9   production_countries                 383995 non-null object
10  release_date                         590327 non-null datetime64[ns]
11  revenue_musd                         14732 non-null  float64
12  runtime                              485664 non-null float64
13  spoken_languages                     377683 non-null object
14  status                               650255 non-null object
15  tagline                              91771 non-null  object
16  title                                650244 non-null object
17  vote_average                         259878 non-null float64
18  vote_count                           650255 non-null int64
19  year                                 650255 non-null object
dtypes: datetime64[ns](1), float64(5), int64(2), object(12)
memory usage: 104.2+ MB

```

```
In [31]: df.isna().sum()
```

```

Out[31]: belongs_to_collection    634079
budget_musd                    619191
genres                          215618
id                               0
original_language               0
overview                       115178
popularity                      0
poster_path                     211457
production_companies            371837
production_countries            266260
release_date                    59928
revenue_musd                    635523
runtime                         164591
spoken_languages                272572
status                           0
tagline                         558484
title                           11
vote_average                     390377
vote_count                       0
year                             0
dtype: int64

```


Removing rows with no "title" values

```
In [32]: df.dropna(subset=['title'], inplace=True)
```

Removing rows that have less than 10 non-null values in total

```
In [33]: df.notna().sum(axis=1).hist(figsize=(8,5), bins=10)
```

```
Out[33]: <AxesSubplot:>
```



```
In [34]: df.notna().sum(axis=1).value_counts()
```

```
Out[34]: 14    84147
         16    81079
         15    80786
         13    77539
         12    74197
         11    72882
         10    60670
          9    42252
         17    36637
          8    18631
         18     9407
         19     5834
          7     4615
         20     1568
dtype: int64
```

```
In [35]: df = df.dropna(thresh=10).reset_index()
```

```
In [36]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 584746 entries, 0 to 584745
Data columns (total 21 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   index                                584746 non-null  int64
 1   belongs_to_collection                16130 non-null   object
 2   budget_musd                         30891 non-null   float64
 3   genres                              430520 non-null  object
 4   id                                   584746 non-null  int64
 5   original_language                   584746 non-null  object
 6   overview                            502425 non-null  object
 7   popularity                          584746 non-null  float64
 8   poster_path                        423837 non-null  object
 9   production_companies                278019 non-null  object
10  production_countries                382146 non-null  object
11  release_date                        557774 non-null  datetime64[ns]
12  revenue_musd                       14696 non-null   float64
13  runtime                             475636 non-null  float64
14  spoken_languages                    375259 non-null  object
15  status                              584746 non-null  object
16  tagline                             91436 non-null   object
17  title                               584746 non-null  object
18  vote_average                        256255 non-null  float64
19  vote_count                          584746 non-null  int64
20  year                                584746 non-null  object
dtypes: datetime64[ns](1), float64(5), int64(3), object(12)
memory usage: 93.7+ MB

```

Final Cleaning Steps

All rows that contain movies that are not released will be dropped, since we are only interested in analyzing released movies.

```
In [37]: df.status.value_counts()
```

```

Out[37]: Released          578044
In Production      2494
Planned            2159
Post Production    1849
Rumored            117
Canceled           83
Name: status, dtype: int64

```

```
In [38]: df = df.loc[df.status == 'Released'].copy()
```


```
In [39]: df
```

Out[39]:

	index	belongs_to_collection	budget_musd	genres	id	original_langu
	0	0	Blondie Collection	NaN	Comedy	3924
	1	1	NaN	NaN	Adventure	6124
	2	2	NaN	NaN	Drama Romance	8773
	3	3	New World Disorder	NaN	NaN	25449
	4	4	NaN	NaN	Family	31975

584741	650250	NaN	NaN	Drama Romance	890928	
584742	650251	NaN	NaN	Comedy	890929	
584743	650252	NaN	NaN	Comedy Fantasy	890930	
584744	650253	NaN	NaN	Family Comedy Romance	890931	
584745	650254	NaN	NaN	Documentary	890932	

578044 rows × 21 columns



In [40]:

df.drop('status', axis=1, inplace=True)

Concatenating the poster_path column values into full URL format and creating a HTML image tag column

In [41]:

df.poster_path[0:4]


```
Out[41]: 0    /o6UMTE2LzQdIKVxRnFECpmtQjsJ.jpg
1    /6xUbUCvndklbGVYiljHr34NTxSl.jpg
2    /aup2QCYCsyEeQfpboXy0f4uj8aE.jpg
3    /okQY6jVmRU19CUbYPUZC77K3XBE.jpg
Name: poster_path, dtype: object
```

```
In [42]: def html_poster_path(value, return_html=False):
        if type(value) == str:
            if return_html == False:
                return f"https://image.tmdb.org/t/p/w500{value}"
            else:
                return f"<img src='https://image.tmdb.org/t/p/w500{value}' style='
        else:
            return np.nan
```

```
In [43]: df['html'] = df['poster_path'].apply(lambda x: html_poster_path(x, True))
df['poster_path'] = df['poster_path'].apply(lambda x: html_poster_path(x))
```

```
In [44]: df[['html', 'poster_path']].head()
```

```
Out[44]:
```

	html	poster_path
0	<img src='https://image.tmdb.org/t/p/w500/o6UM...	https://image.tmdb.org/t/p/w500/o6UMTE2LzQdIKV...
1	<img src='https://image.tmdb.org/t/p/w500/6xUb...	https://image.tmdb.org/t/p/w500/6xUbUCvndklbGV...
2	<img src='https://image.tmdb.org/t/p/w500/aup2...	https://image.tmdb.org/t/p/w500/aup2QCYCsyEeQf...
3	<img src='https://image.tmdb.org/t/p/w500/okQY...	https://image.tmdb.org/t/p/w500/okQY6jVmRU19CU...
4	<img src='https://image.tmdb.org/t/p/w500/qKWc...	https://image.tmdb.org/t/p/w500/qKWcCmvGr4g0dg...

Next, credits and cast information will be added to the DataFrame from the previously downloaded dataset.

```
In [3]: credits = pd.read_csv('credits.csv')
credits
```

Out[3]: 0

```

0   {'id': 3924, 'cast': [{'adult': False, 'gender...
1   {'id': 6124, 'cast': [{'adult': False, 'gender...
2   {'id': 8773, 'cast': [{'adult': False, 'gender...
3   {'id': 25449, 'cast': [{'adult': False, 'gende...
4   {'id': 31975, 'cast': [], 'crew': []}
...
650241 {'id': 890928, 'cast': [], 'crew': []}
650242 {'id': 890929, 'cast': [], 'crew': []}
650243 {'id': 890930, 'cast': [{'adult': False, 'gend...
650244 {'id': 890931, 'cast': [{'adult': False, 'gend...
650245 {'id': 890932, 'cast': [{'adult': False, 'gend...

```

650246 rows × 1 columns

In [4]: `credits.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 650246 entries, 0 to 650245
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0    0      650246 non-null  object
dtypes: object(1)
memory usage: 5.0+ MB

```

Evaluating each value using `literal_eval`

In [46]: `credits['0'] = literal_eval(credits['0'])`

In [47]: `credits = pd.json_normalize(credits['0'])`
`credits.head()`

Out[47]:

	id	cast	crew
0	3924	[{'adult': False, 'gender': 1, 'id': 34178, 'k...	[{'adult': False, 'gender': 0, 'id': 34170, 'k...
1	6124	[{'adult': False, 'gender': 2, 'id': 48038, 'k...	[{'adult': False, 'gender': 2, 'id': 2902, 'kn...
2	8773	[{'adult': False, 'gender': 2, 'id': 1653, 'kn...	[{'adult': False, 'gender': 2, 'id': 1650, 'kn...
3	25449	[{'adult': False, 'gender': 0, 'id': 84130, 'k...	[{'adult': False, 'gender': 0, 'id': 112786, '...
4	31975	[]	[]

Replacing empty list entries with NaN

In [30]: `for column in ['cast', 'crew']:`
`credits[column] = credits[column].apply(lambda x: np.nan if len(x) == 0 el`

```
In [31]: credits.isnull().sum()
```

```
Out[31]: id          0
         cast    217840
         crew    116372
         dtype: int64
```

Dropping rows with less than 2 entries

```
In [32]: credits.dropna(thresh=2, inplace=True)
```

Extracting the names of both cast and crew columns for each row

```
In [33]: credits['cast_names'] = credits['cast'].apply(lambda x: '|'.join([i['name'] for i in x]))
         credits['crew_names'] = credits['crew'].apply(lambda x: '|'.join([i['name'] for i in x]))
```

Getting cast and crew size for each movie.

```
In [34]: credits['cast_size'] = credits['cast'].apply(lambda x: len(x) if type(x) != float else 0)
         credits['crew_size'] = credits['crew'].apply(lambda x: len(x) if type(x) != float else 0)
```

Extracting the director name for each movie

```
In [35]: def get_director(x):
         if type(x) != float:
             for i in x:
                 if i['job'] == 'Director':
                     return i['name']
             return np.nan
```

```
In [36]: credits['director'] = credits['crew'].apply(get_director)
```

```
In [37]: credits['director'].value_counts(dropna=False)
```

```
Out[37]: NaN          46510
         Dave Fleischer    494
         D. W. Griffith    417
         Kevin Dunn       373
         Stan Brakhage     358
         ...
         Cecily Fay        1
         Raaghav           1
         Yann Sông Le Van Ho 1
         Kim Hyun-joon     1
         Junca Avilés      1
         Name: director, Length: 186778, dtype: int64
```

Removing duplicate entries

```
In [38]: credits.drop_duplicates(subset='id', inplace=True)
```

Joining Both DataFrames

```
In [81]: combined_df = df.merge(credits[['id', 'cast_names', 'crew_names', 'director']])
```

```
In [82]: combined_df.head()
```

```
Out[82]:
```

	index	belongs_to_collection	budget_musd	genres	id	original_language	overview
0	0	Blondie Collection	NaN	Comedy	3924	en	Blondie and Dagwood are about to celebrate the...
1	1	NaN	NaN	Adventure	6124	de	Der Mann ohne Namen is a German adventure movi...
2	2	NaN	NaN	Drama Romance	8773	fr	Love at Twenty unites five directors from five...
3	3	New World Disorder	NaN	NaN	25449	en	Gee Atherton ripping the Worlds course the day...
4	4	NaN	NaN	Family	31975	en	Elmo is making a very, very super special surp...

5 rows × 24 columns

Saving the final DataFrame into a csv file.

```
In [83]: combined_df.to_csv('movies_complete.csv', index = False)
```

In Part 2 of this project, we will try and gain some insight from the cleaned data about what types of movies are more popular, the

average budget and revenue of the most successful movies, what genres are more popular, and many more.